

# Let's master Hibernate!

Michał Żmuda  
Piotr Turek

# Agenda

- Transactions
  - Why do you need them - example
  - ACID
  - Building complex bussiness logic - how to use propagation
  - Transaction Manager - go declarative!
  - Propagation levels

# Agenda

- Review of JPA compatibility
  - What is it about Hibernate vs JPA - a bit of history
  - Interoperability and replacements
  - Hibernate vs JPA - which one to choose?

# Agenda

- Things to consider designing ORM software
  - Eager fetching performance
    - case showing disastrous performance
    - actual/typical overhead
  - Lazy initialization pitfalls
  - Caching - When, How, Why, plus examples
  - Cascades - how not to hurt yourself
  - Golden rules of hibernate based orm layer design

# Transactions - none

```
OrdersDao ordersDao = (OrdersDao) ctx.getBean("ordersDao");
OrderDetailsDao orderDetailsDao = (OrderDetailsDao) ctx.getBean("orderDetailsDao");
OrdersEntityFactory ordersEntityFactory = (
    OrdersEntityFactory) ctx.getBean("ordersEntityFactory");

OrdersEntity randomOrder = ordersEntityFactory.createRandomOrder();
ordersDao.create(randomOrder);

OrderDetailsEntity detailsEntity = new OrderDetailsEntity();
detailsEntity.setOrderid(randomOrder.getOrderid());
setOtherOrderDeatilsFields(orderDetailsDao, detailsEntity);
orderDetailsDao.create(detailsEntity); //what happens when it fails?
```

- imagine that creating OrderDetails fails  
(exercise: make it fail! tip: violate constraint)
- we have only Order record but we do not have any data what send anf for how much
- how to recover from such failure
- solution: perform both creates or any at all

# Transactions

```
// session created by openSession() will be used by dao's - see GenericDao!  
Transaction transaction = sessionFactory.openSession().beginTransaction();  
OrdersEntity randomOrder = ordersEntityFactory.createRandomOrder();  
ordersDao.create(randomOrder);  
  
OrderDetailsEntity detailsEntity = new OrderDetailsEntity();  
detailsEntity.setOrderid(randomOrder.getOrderid());  
setOtherOrderDeatilsFields(orderDetailsDao, detailsEntity);  
orderDetailsDao.create(detailsEntity);  
transaction.commit();
```

- now sequence of creates is protected by orm transaction mechanism
- transactions provide easy to use yet powerful help in implementing business logic
- they are analogous to database transactions
- they follow ACID acronym

# Transactions

- are limited to one session  
see session opened in example and how DAOs are designed
- may be used by Transaction Manager  
discussed later - manual management better serves examples  
TM used by annotating methods @Transactional
- manually handled by struct like:

```
Transaction transaction = null;
try {
    transaction = session.beginTransaction();
    //operations
    transaction.commit();
} catch (Exception e) {
    if (transaction != null) {
        transaction.rollback();
    }
}
```

# ACID - Atomicity

- previous example shows need of atomicity
- grouping set of changes in transaction guarantees that they all be executed or any of them at all
- custom fallback policies may be handled with savepoints
- savepoint is used to group changes into atomic groups inside an transaction
- savepoints are provided by Transaction Manager which will be discussed later
- question: do you see any application of savepoints in our northwind example?



# ACID - Consistency

- imagine that in db exists trigger that order has at most 10 order detail records
- when it would be violated inside a transaction, it would be all rolled back
- this guarantees consistency - all data after transaction fulfill validation rules

optional exercise: try adding such trigger to the db and prepare runner violating it inside transaction

# ACID - Isolation

```
executor.submit((Runnable) () -> {  
    try {  
        Thread.currentThread().sleep(4000);  
        printTotalValueStatistic(orderDetailsDao,ordersDao);  
    } catch (InterruptedException e) {  
        e.printStackTrace();  
    }  
});  
printTotalValueStatistic(orderDetailsDao, ordersDao);  
OrdersEntity randomOrder = ordersEntityFactory.createRandomOrder();  
ordersDao.create(randomOrder);  
Thread.currentThread().sleep(6000); //simulating time window  
OrderDetailsEntity detailsEntity = new OrderDetailsEntity();  
detailsEntity.setOrderid(randomOrder.getOrderid());  
setOtherOrderDeatilsFields(orderDetailsDao, detailsEntity);  
orderDetailsDao.create(detailsEntity);  
printTotalValueStatistic(orderDetailsDao,ordersDao);
```

- do you see what might occur?
- do you believe that without sleep() such situation won't never happen?

# ACID - Isolation

```
Total value of 1.120000137685959E13 is in 858 orders  
Total value of 1.120000137685959E13 is in 859 orders  
Total value of 1.680000137685959E13 is in 859 orders
```

- middle result was generated when order was already create but it's details not
- when statistics and order addition executed in different transactions it is guaranteed that changes from one transaction won't be visible to another until it's completed  
(this is overall description - isolation levels later)
- question: what middle result will look like if transactions were used?
- exercise: introduce transactions

# ACID - Isolation

for better isolation we pay with performance - choose of isolation level is given

- READ UNCOMMITTED - no protection, see 'dirty reads'
  - READ COMMITTED - when reading row you might get different values due update
    - see 'non-repeatable reads'
  - REPEATABLE READ - when selecting range of rows you might get different result sets
    - see 'phantom read'
  - SERIALIZABLE - whole tables are locked - safest
- isolation levels may be set declaratively with Transaction Manager or programatically with TransactionTemplates
- typically we annotate method to be handled by TM with annotation where we can set isolation level

```
@Transactional(isolation = Isolation.SERIALIZABLE)
```

question: which level is required to make our example work?

# ACID - Durability

- durability means that after transaction completion all changes made within it are now persisted
- this is core requirement for persistence layer and won't be discussed further



# Transaction Manager

- allows to declaratively use transactions
- configured once in applicationContext.xml

```
<tx:annotation-driven/>  
<bean id="transactionManager" class=  
    "org.springframework.orm.hibernate4.HibernateTransactionManager">  
    <property name="sessionFactory" ref="sessionFactory"/>  
</bean>
```

- by annotating method with @Transactional we are guaranteed that its instructions will be executed in transaction
- it may be new transaction or one already opened in method invoking our method
- this behaviour is controlled by propagation levels

# TM - propagation levels

```
@Service
public class OrdersValueService {
    @Transactional(propagation = Propagation.NESTED)
    public void printTotalValueStatistic() {
```

```
@Transactional
public void addNewOrder() throws InterruptedException {
    OrdersEntity randomOrder = ordersEntityFactory.createRandomOrder();
    ordersDao.create(randomOrder);
    ordersValueService.printTotalValueStatistic(); //during
    OrderDetailsEntity detailsEntity = new OrderDetailsEntity();
    detailsEntity.setOrderid(randomOrder.getOrderid());
    setOtherOrderDeatilsFields(orderDetailsDao, detailsEntity);
    orderDetailsDao.create(detailsEntity);
}
```

```
public static void main(String[] args) throws InterruptedException {
    ClassPathXmlApplicationContext ctx = new ClassPathXmlApplicationContext("META-INF/a
    ((OrdersValueService)ctx.getBean("ordersValueService")).printTotalValueStatistic();
    ((PropagationLevelsRunner) ctx.getBean("propagationLevelsRunner")).run();
    ((OrdersValueService)ctx.getBean("ordersValueService")).printTotalValueStatistic();
}
```

# TM - propagation levels

- execution of example yields

```
Total value of 5.600001376863591E12 is in 861 orders  
Total value of 5.600001376863591E12 is in 862 orders  
Total value of 1.120000137686359E13 is in 862 orders
```

- this result is similar to one we would get without transactions but we use them (each method is using TM thanks to annotation)
- see propagation level - this particular used in the example causes that actually we have only one transaction!



# TM - propagation levels

- MANDATORY Support a current transaction, throw an exception if none exists.
- NESTED Execute within a nested transaction if a current transaction exists, behave like PROPAGATION\_REQUIRED else.
- NEVER Execute non-transactionally, throw an exception if a transaction exists.
- NOT\_SUPPORTED Execute non-transactionally, suspend the current transaction if one exists.
- REQUIRED Support a current transaction, create a new one if none exists.
- REQUIRES\_NEW Create a new transaction, suspend the current transaction if one exists.
- SUPPORTS Support a current transaction, execute non-transactionally if none exists.

question: which levels would fix our example? try it out!

# Hibernate vs JPA

*What JPA is?*

- a specification

*What JPA isn't?*

- an implementation



*What Hibernate is?*

- an implementation

*What Hibernate isn't?*

- a specification



Hibernate provides a superset of JPA functionality!

... just like a musician can play more than one song or symphony

# Hibernate vs JPA

## Hibernate vs JPA

*equals*

SessionFactory vs EntityManagerFactory

*equals*

implementation coupling vs implementation independence

*equals*

powerfull features vs limited functionality

# Hibernate vs JPA

Hibernate	JPA
SessionFactory	EntityManagerFactory
Session	EntityManager
sessionFactory.getCurrentSession(). [method]()	entityManager.[method]()
saveOrUpdate()	persist()
Query.setInteger/String/Entity()	Query.setParameter()
list()	getResultList()
uniqueResult()	getSingleResult()
uniqueResult() returns null	getSingleResult() throws NoResultException
CriteriaQueries – yes	CriteriaQueries – no

- You can unwrap EntityManager to obtain Session
  - entityManager.getDelegate()
  - entityManager.unwrap(Session.class)

# Hibernate vs JPA - which to choose

Advantages of using JPA:

- designing and coding to API
- standardization
- increased portability (!)

*Maybe one day you will want to change the provider to something different than Hibernate?*

... quite likely however, that will never ever happen

Therefore ...

# Hibernate vs JPA - which to choose

... considering:

- much better exception translation
- more power, more features

Use Hibernate's Session and Session Factory

Good compromise is to use JPA annotations with Hibernate Session

*... but, most probably you do that anyway already*



# Eager fetch performance

simple example of performance:

(required service can be found in project or implemented more elegant on your own)

```
Collection<ProductsEntity> products = service.generateProducts(productsDao, 10000);
OrdersEntity randomOrder = prepareOrdersEntityWithOrderDetails(ordersDao, products,
service.createOrderDetailsFromOrderEntity(randomOrder));
long time = System.nanoTime();
System.out.println(ordersDao.get(randomOrder.getOrderid()).getOrderdate());
System.out.println("Elapsed " + (System.nanoTime() - time) + "ns");
```

```
Hibernate: select ordersenti0_.orderid as orderid1_7_4_, ordersenti0_
2014-05-13
Elapsed 25468826
```

by default hibernate lazy initialisation, change to eager fetching

```
@OneToMany(mappedBy = "ordersByOrderid", fetch = FetchType.EAGER)
public Collection<OrderDetailsEntity> getOrderDetailsesByOrderid() {
    return orderDetailsesByOrderid;
}
```

```
Hibernate: select ordersenti0_.orderid as orderid1_7_8_, ordersenti0_.customerid a
2014-05-13
Elapsed 2454467513ns
```

# Eager fetch performance

- that's 100 times slower!
- moreover it's over 2s to read simple value!
- and it could be more than related records 10000...

- don't want ever to use eager fetchnig?

you may be wrong - it prooves quite usefull when you have relations @OneToOne or @ManyToOne and it's used by default with such relations by default by JPA

question: if forbidden to change fetching type, how you would tweak reading single value from the db?



# Lazy initialization hell

Eager fetching is bad.

Lazy fetching is treacherous.

```
@OneToMany(mappedBy = "productsByProductid", fetch = FetchType.LAZY)
public Collection<OrderDetailsEntity> getOrderDetailsesByProductid() { return orderDetailsesByProductid; }
```

```
private void run() {
    final ProductsEntity product = lazyInitHellService.extractProduct();
    displayOrderDetails(product);
}
```

```
@Transactional
public ProductsEntity extractProduct() {
    return productsDao.getAll().get(0);
}
```

```
private void displayOrderDetails(ProductsEntity product) {
    final Collection<OrderDetailsEntity> orderDetails = product.getOrderDetailsesByProductid();
    final OrderDetailsEntity firstDetail = orderDetails.iterator().next();
    System.out.println("Discount for a product: " + firstDetail.getDiscount());
}
```

Simple. What can possibly go wrong... (try it out)

# Lazy initialization hell

... well,

```
Exception in thread "main" org.hibernate.LazyInitializationException: failed to lazily initialize a
collection of role: pl.agh.turek.bazy.hibernate.model.ProductsEntity.orderDetailsesByProductid, could
not initialize proxy - no Session <5 internal calls>
    at pl.agh.turek.bazy.hibernate.runners.LazyInitHellRunner.displayOrderDetails(LazyInitHellRunner
.java:32)
    at pl.agh.turek.bazy.hibernate.runners.LazyInitHellRunner.run(LazyInitHellRunner.java:27)
    at pl.agh.turek.bazy.hibernate.runners.LazyInitHellRunner.main(LazyInitHellRunner.java:22)
    <5 internal calls>
```

## Explanation:

- Hibernate returns a HibernateProxy, not a real ProductEntity
- HibernateProxy instance belongs to a Session that created it
- When transaction ends, session gets closed
- When accessing lazy collection, Hibernate tries to read it lazily
  - but the session had already been closed (!)

# Lazy initialization hell

```
@Transactional
public ProductsEntity extractProductWithOrderDetails() {
    final ProductsEntity product = extractProduct();
    System.out.println("... Product lazy proxy retrieved ...");
    product.getOrderDetailsesByProductid().size();
    return product;
}
```

We need to “unpack” the collection while the session is open

- we can do that f.e. by calling size() method on it

```
... Product lazy proxy retrieved ...
```

```
Hibernate: select orderdetail0_.productid as product1_8_6_, orderdetail0_.productid as product1_6_6_, orderdetail0_.orderid as orderdetail0_.discount as discount3_6_5_, orderdetail0_.quantity as quantity4_6_5_, orderdetail0_.unitprice as unitpric5_6_5_, ordersentil_.employeeid as employee3_7_0_, ordersentil_.freight as freight4_7_0_, ordersentil_.orderdate as orderdat5_7_0_, ordersentil_.shipcity as shipcity8_7_0_, ordersentil_.shipcountry as shipcoun9_7_0_, ordersentil_.shipname as shipnam10_7_0_, ordersentil_.shippostalcode as shippos12_7_0_, ordersentil_.shipregion as shipreg13_7_0_, customerse2_.customerid as customerse2_.companyname as companyn4_3_1_, customerse2_.contactname as contactn5_3_1_, customerse2_.contacttitle as contact customerse2_.phone as phone9_3_1_, customerse2_.postalcode as postalcl10_3_1_, customerse2_.region as region11_3_1_, employee as birthdat3_4_2_, employeeese3_.city as city4_4_2_, employeeese3_.country as country5_4_2_, employeeese3_.reportsto as reports employeeese3_.hiredate as hiredate8_4_2_, employeeese3_.homephone as homephon9_4_2_, employeeese3_.lastname as lastnam10_4_2_, as photopal3_4_2_, employeeese3_.postalcode as postalcl14_4_2_, employeeese3_.region as region15_4_2_, employeeese3_.title as ti employeeese4_.employeeid as employee1_4_3_, employeeese4_.address as address2_4_3_, employeeese4_.birthdate as birthdat3_4_3_, as reports16_4_3_, employeeese4_.extension as extensio6_4_3_, employeeese4_.firstname as firstnam7_4_3_, employeeese4_.hiredate
```

HibernateProxy is lazily fetching data from database.

Problem Solved

:)

# Caching - when, how and why

## Problem:

- There exist entities that have much more reads than writes
  - think of an online shop and ProductsEntity
  - details rarely updated (maybe even never)
  - thousands of clients read data, every minute
- Going to database for that data each time is an overkill
  - and can literally kill your database (and app as well)

## Solution:

- Add some caching ...
  - which can give some pretty amazing results

# Caching - configuration

Hibernate properties:

```
<prop key="hibernate.cache.use_second_level_cache">true</prop>
<prop key="hibernate.cache.provider_class">org.hibernate.testing.cache.CachingRegionFactory</prop>
<prop key="hibernate.cache.region.factory_class">org.hibernate.testing.cache.CachingRegionFactory</prop>
```

- enable second level caching
- set up caching region provider
  - Hibernate's ConcurrentHashMap-based one works ok
- add missing dependency

```
<dependency>
  <groupId>org.hibernate</groupId>
  <artifactId>hibernate-testing</artifactId>
  <version>${hibernate-version}</version>
</dependency>
```



# Caching - setting up caches

Selected entities need to be marked as cacheable.

```
@Entity
@Cacheable
@Cache(usage = CacheConcurrencyStrategy.READ_WRITE)
@Table(name = "products", schema = "public", catalog = "northwind")
public class ProductsEntity {
```

Two annotations:

- @Cacheable - marker annotation
- @Cache(...)
  - region(), *The cache region. This attribute is optional, and defaults to the fully-qualified class name of the class, or the fully-qualified role name of the collection.*

# Caching - setting up caches

- @Cache(...)
  - include(), *Whether or not to include all properties.*
  - usage(),
    - *read-only - entity has read-only access*
    - *read-write - support for writing, requires locking*
    - *nonstrict-read-write - allows writing without locking (avoid)*
    - *transactional - full transactional support*

Read-only caches are to be preferred!

# Caching - deadly trap

Lets try an excercise:

- Configure caching for ProductsEntity
- Run LazyInitHellRunner
- Modify data manually in a database (in Products table)
  - For example remove the first row
- Run the runner again

Result:

- Data integrity goes to hell(!)
- which can and most probably will, break your business logic
- ... and cost you money



# Caching - deadly trap (how to avoid)

## Conclusion:

- **NEVER, EVER** modify data externally, when using caching
  - no manual inserts/updates/deletes
  - no external systems modifying db state in the background
- **IF** for some reason such events can occur:
  - **notify** your system about this fact
  - **evict** the caches:

```
@Override
public void purgeAllCaches() {
    final Cache cache = sessionFactory.getCache();
    cache.evictEntityRegions();
    cache.evictCollectionRegions();
    cache.evictDefaultQueryRegion();
    cache.evictQueryRegions();
}
```

# Cascades - how not to hurt yourself

Consider following code: (try it out - try to explain results)

```
private void run() {  
    final ProductsEntity productsEntity = lazyInitHellService.extractProductWithOrderDetails();  
    final OrderDetailsEntityPK pk = cascadeTypeService.modifyFirstOrderDetailsDiscount(productsEntity);  
    cascadeTypeService.verifyOrderDetailsDiscount(pk);  
}
```

```
@Transactional  
public OrderDetailsEntityPK modifyFirstOrderDetailsDiscount(ProductsEntity product) {  
    System.out.println("Updating discount of first order detail of product id: " + product.getProductid());  
    final Collection<OrderDetailsEntity> orderDetails = product.getOrderDetailsesByProductid();  
    final OrderDetailsEntity firstOrderDetail = orderDetails.iterator().next();  
    final OrderDetailsEntityPK pk = getOrdersPK(firstOrderDetail);  
    final double discount = firstOrderDetail.getDiscount();  
    System.out.println("Will try to fetch order detail id:" + pk + " with discount "  
        + discount);  
    firstOrderDetail.setDiscount(discount + 0.1);  
    System.out.println("New discount set to: " + firstOrderDetail.getDiscount());  
    productsDao.update(product);  
    return pk;  
}
```

```
@Transactional  
public void verifyOrderDetailsDiscount(OrderDetailsEntityPK pk) {  
    final OrderDetailsEntity refetchedOrderDetail = orderDetailsDao.get(pk);  
    System.out.println("Order details id " + pk + " has discount " + refetchedOrderDetail.getDiscount());  
}
```

# Cascades - how not to hurt yourself

The code:

- fetches a product
- adds 0.1 to the discount of the first order
- updates the product
- verifies the value of discount

When you run it without any changes to the model:

```
Updating discount of first order detail of product id: 5
Will try to fetch order detail id:OrderDetailsEntityPK{orderid=10258, productid=5} with discount 0.200000003
New discount set to: 0.300000003
Hibernate: update northwind.public.products set categoryid=?, discontinued=?, productname=?, quantityperunit=?, reord
Order details id OrderDetailsEntityPK{orderid=10258, productid=5} has discount 0.200000003
```

As you can see, even though the product has been updated, the discount stayed unchanged.

That's because, *by default, no operations are “propagated” to relations*

# Cascades - how not to hurt yourself

**cascadeType** attribute:

- ALL - Cascade all operations
- DETACH - Cascade detach operation
- MERGE - Cascade merge operation
- PERSIST - Cascade persist operation
- REFRESH - Cascade refresh operation
- REMOVE - Cascade remove operation

Lets use it on our ProductEntity:

```
@OneToMany(mappedBy = "productsByProductid", fetch = FetchType.LAZY, cascade = {CascadeType.MERGE, CascadeType.PERSIST})  
public Collection<OrderDetailsEntity> getOrderDetailsesByProductid() {
```

**Result:**

```
Updating discount of first order detail of product id: 7  
Will try to fetch order detail id:OrderDetailsEntityPK{orderid=10262, productid=7} with discount 0.0  
New discount set to: 0.1
```

```
Order details id OrderDetailsEntityPK{orderid=10262, productid=7} has discount 0.1
```

*Merge operation has been “propagated” to relations*

# Cascades - how not to hurt yourself

## However:

- Cascading must be used with great care
- ... and great deal of thought

## Otherwise:

- We can literally implode our database with a single `remove()`

## Thankfully:

- Database constraints can work as a safety measure against such scenario

## In general:

- Avoid using *CascadeType.ALL*